

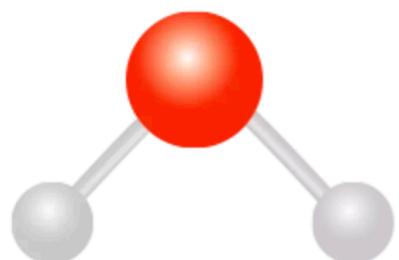
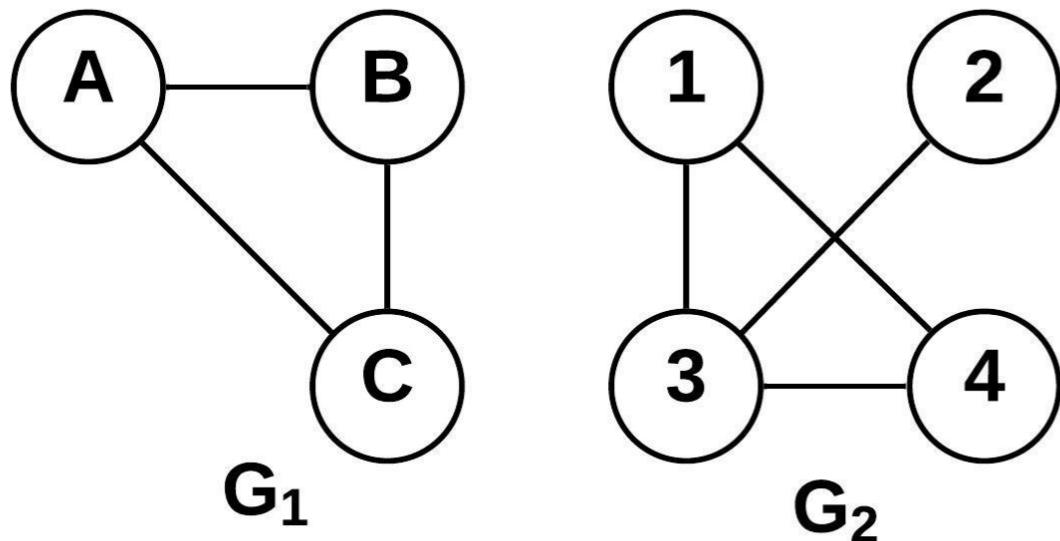
A PARALLEL ALGORITHM FOR SUBGRAPH ISOMORPHISM

Rohan Yadav and Umut A. Acar

Carnegie Mellon University

SPAA 2019

SUBGRAPH ISOMORPHISM BACKGROUND



A large, faint watermark is centered over the binary code matrix. It features the number '10101010' repeated twice in a vertical column. Above this, there is a circular outline containing the binary sequence '1111100011'. Below the main text, there is a horizontal line with a small circle at its center, followed by the binary sequence '101101100011'. The watermark is rendered in a light gray color that blends with the background.

- Is G_1 a subgraph of G_2 ?
 - Involves finding a mapping f of vertices from G_1 to G_2 that preserves edge relationships
 - For each edge (u, v) in G_1 , $(f(u), f(v))$ must be in G_2
 - Subgraph isomorphism is NP-Complete
 - Has applications in
 - pattern recognition
 - biochemical applications
 - graph databases
 - ...

STRATEGY: SEARCH ALGORITHMS

- Explore a state space to find the isomorphism
- A **match** is a tuple of vertices mapping a vertex in G_1 to a vertex in G_2
- A **matching** is a partial isomorphism between G_1 and G_2 , represented by a set of matches
- Matchings grow and shrink by adding and removing matches
- A matching is **consistent** if it preserves edge relationships
- **Algorithm** —

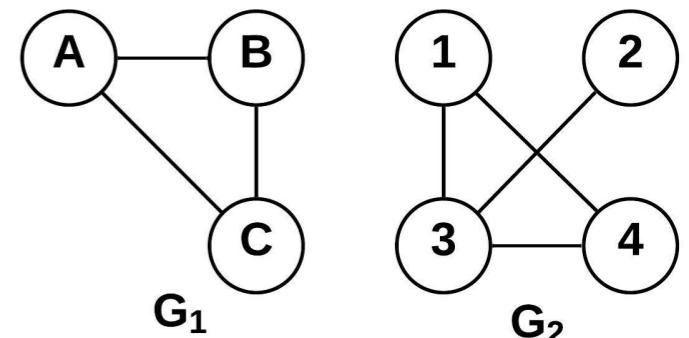
Start with an empty matching M

Loop:

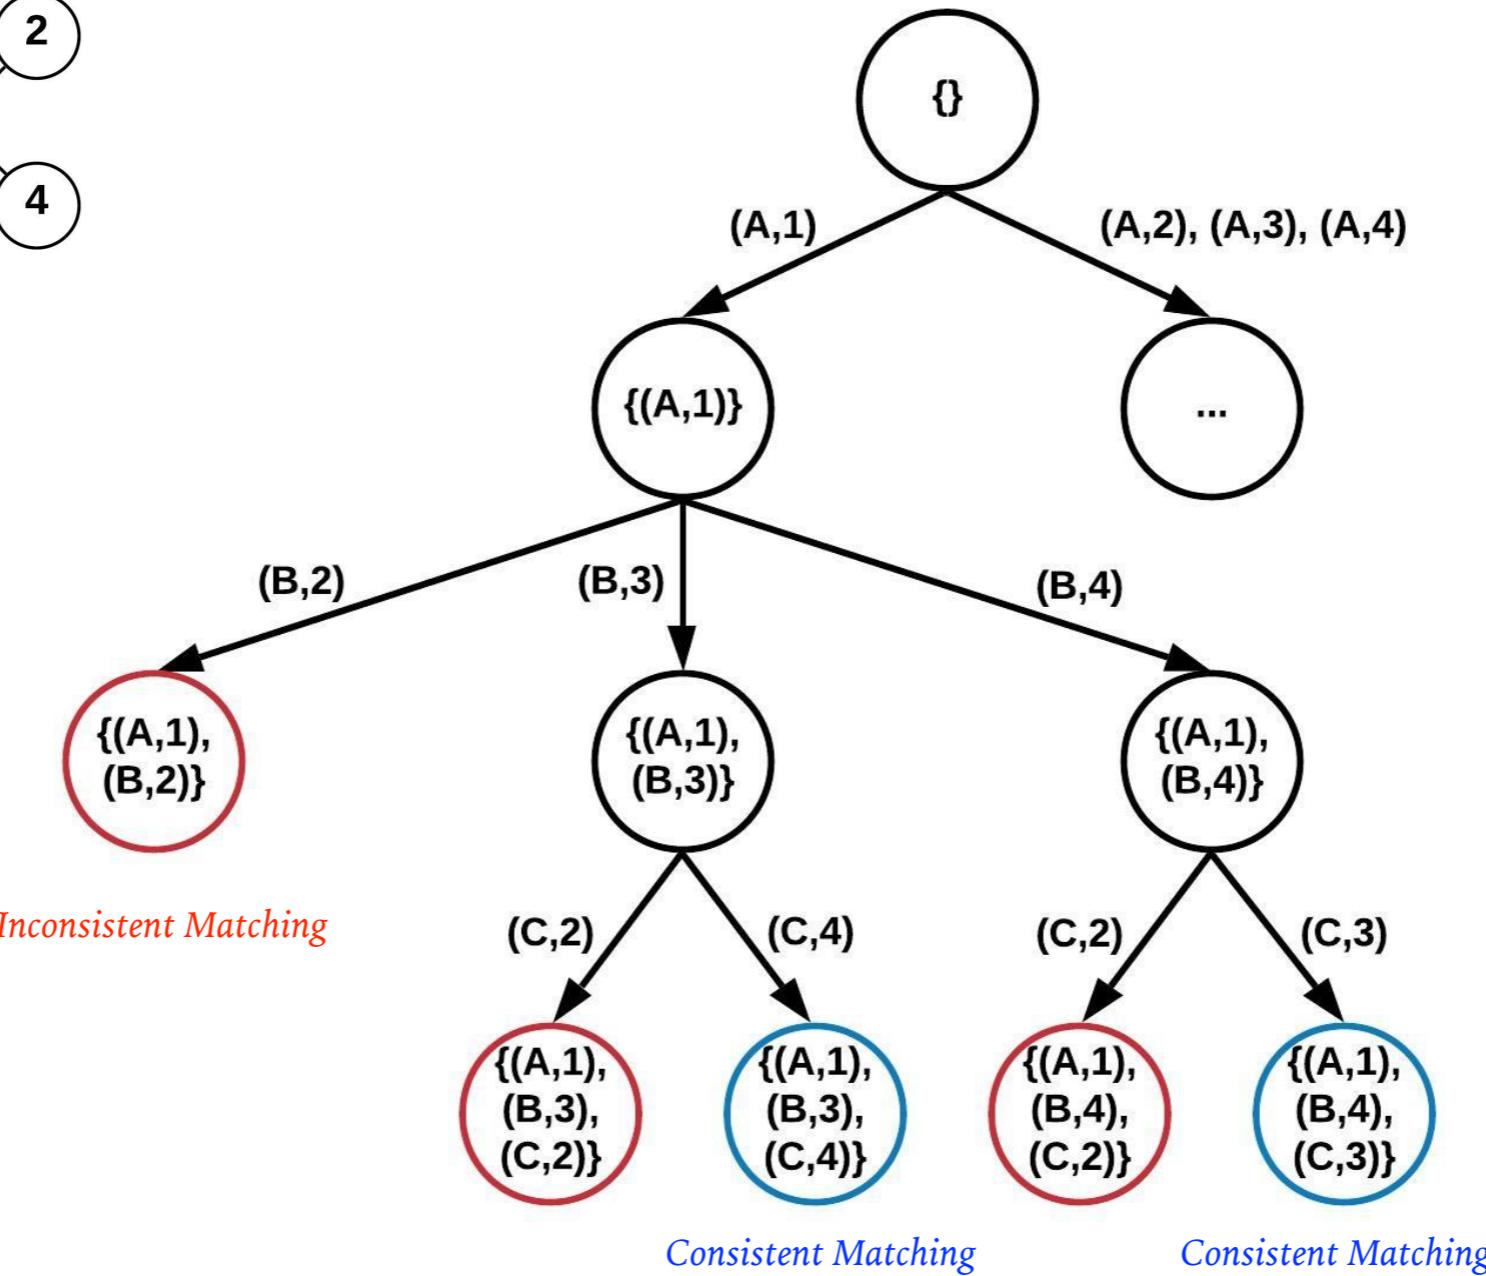
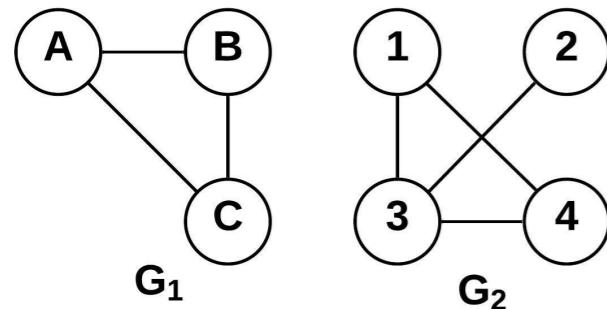
add a new match (u, v) to M such that M is not visited

visit M if M is consistent

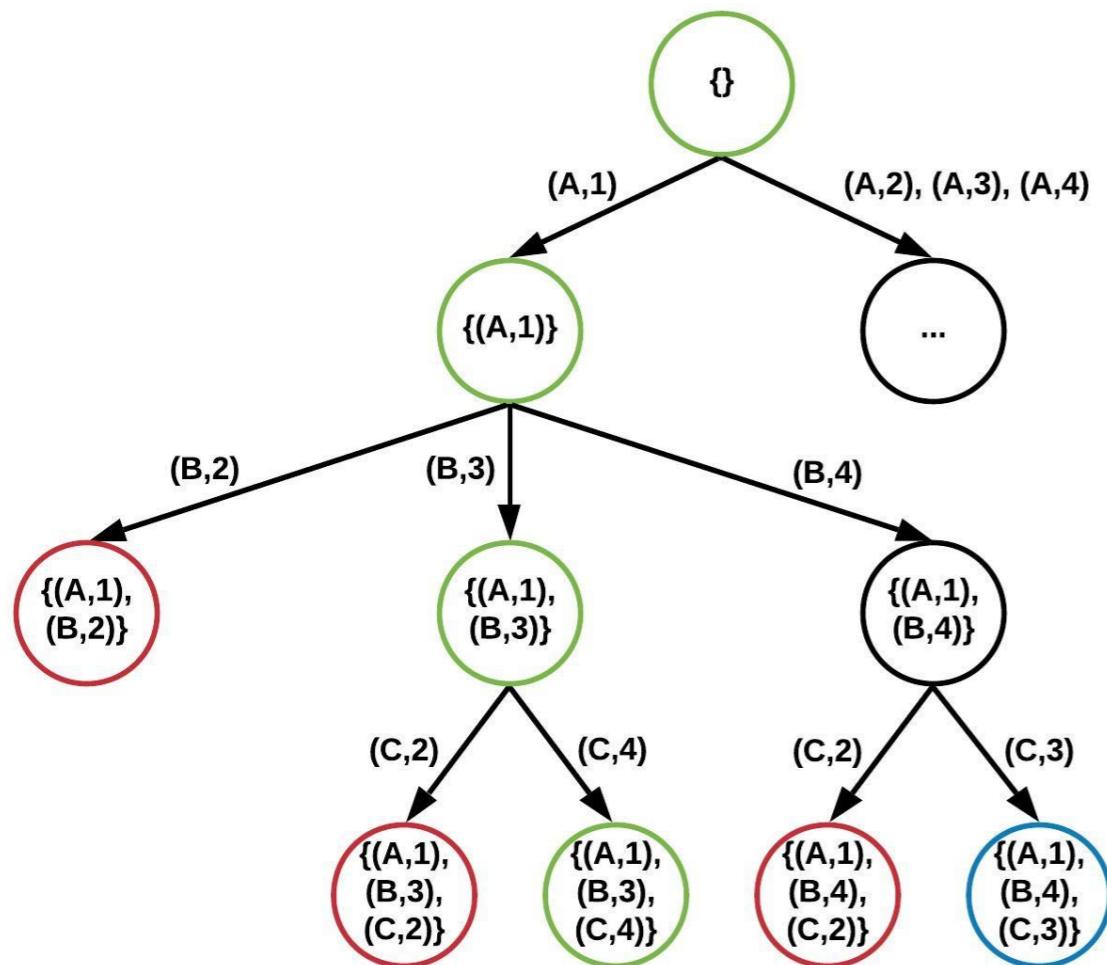
if M is isomorphism then done



SEARCH TREE STRUCTURE



BASIC TREE SEARCH ALGORITHM



```
def search M =  
    for each child edge (u, v) of M:  
        if M + (u, v) is consistent:  
            add (u, v) to M  
        if search M:  
            return true  
        remove (u, v) from M  
    return false
```

- Key idea: efficient backtracking
- Core of known algorithms like VF2, RI [Cordella '04, Bonnici '13]

CHALLENGES WITH PARALLELIZATION

```
def search M =  
    parallel for each child edge  
    (u, v) of M:  
        if M + (u, v) is consistent:  
            add (u, v) to M  
        if search M:  
            return true  
        remove (u, v) from M  
    return false
```

- Idea: try all children in parallel
- Performs poorly — why?
- Requires persistent matching structure
- Highly irregular branches
 - With pruning, search tree is irregular
 - Work is not predictable resulting in fine grained tasks
- Prior work has identified these issues

LAZY PARALLELISM

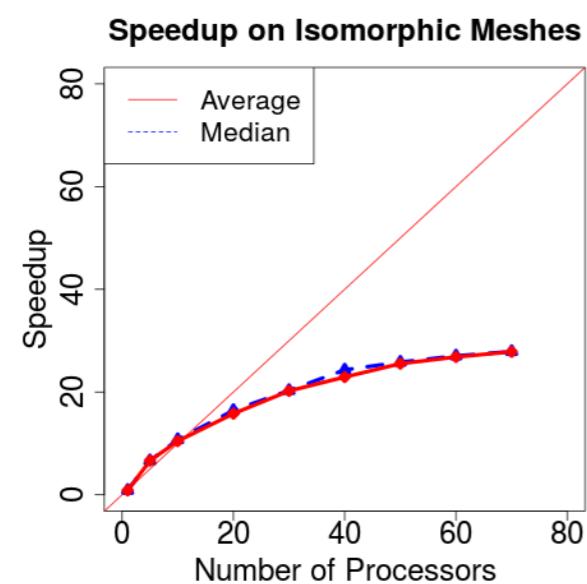
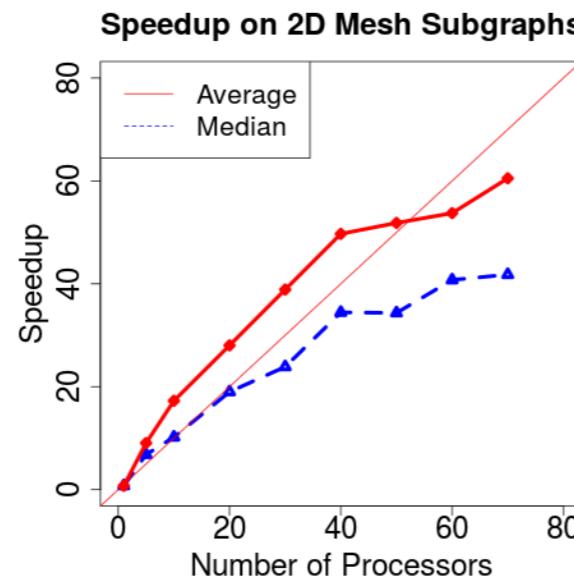
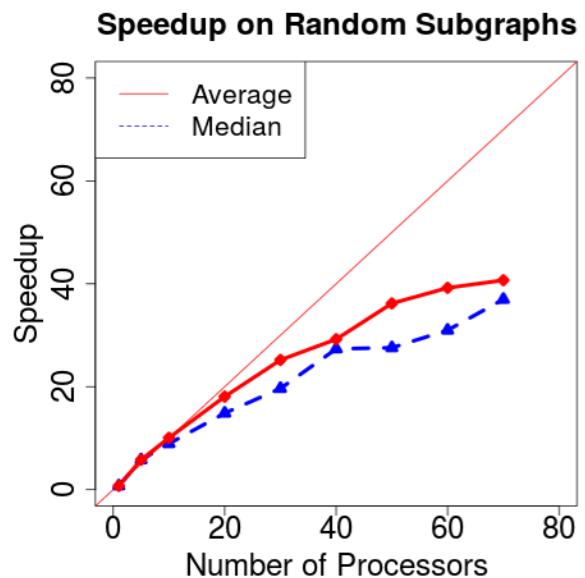
- Goal: maintain efficiency of sequential algorithm
- Create parallelism lazily on-demand
 - p workers
 - Each worker runs (almost) the sequential algorithm
 - But generates parallelism by splitting its work when requested



ALGORITHMS FOR LAZY SPLITTING

- Frontier data structure for representing work of each worker
 - Push and pop edges from frontier to explore
 - Split frontier when requested to share work
- Backtracking from failed searches
 - Frontier structure represents a search path
 - Split operation returns contiguous search paths
 - Frontier invariants make it possible to efficiently implement backtracking
- Scheduling for irregular and unpredictable parallelism
 - Amortization technique: workers share work only after performing enough work to pay for sharing [Acar et al. '15]

PRELIMINARY RESULTS



Subset of preliminary results

- Based on VF2 implementation
- Tested on the database designed by the creators of the VF2 algorithm
- Overhead over VF2 algorithm is between 15%-40%
- Established bounds on single core work efficiency and p -processor space overhead

FUTURE WORK

- Extend the core ideas of our algorithm to more recent algorithms like VF3 [Vincenzo '16]
- Perform a larger experimental evaluation with more data sets and different metrics (aggregate speedup, etc.)
- Explore theory and analysis of lazy splitting style algorithms, prove guarantees about the span of our algorithm